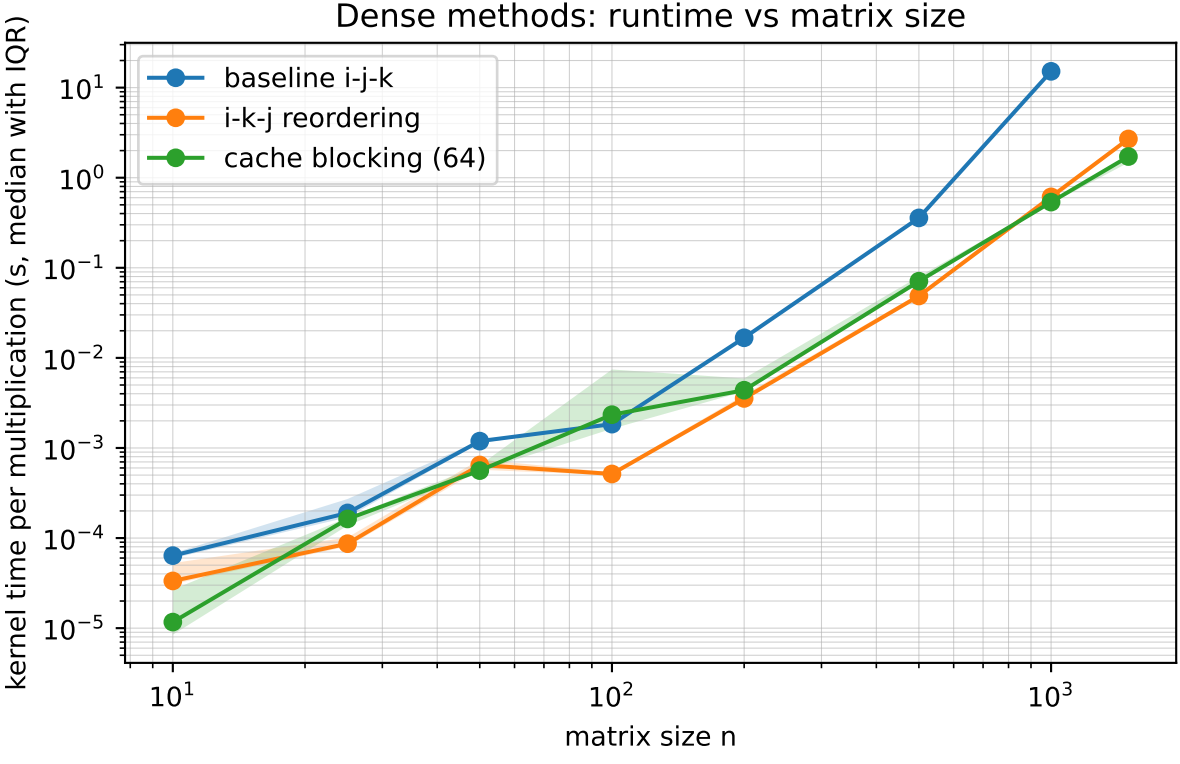
Dense methods: runtime vs matrix size
baseline i-j-k
i-k-j reordering
cache blocking (64)
kernel time per multiplication (s, median with IQR)
matrix size n
10^{1}
10^{0}
10^{-1}
10^{-2}
10^{-3}
10^{-4}
10^{-5}
10^{1}
10^{2}
10^{3}