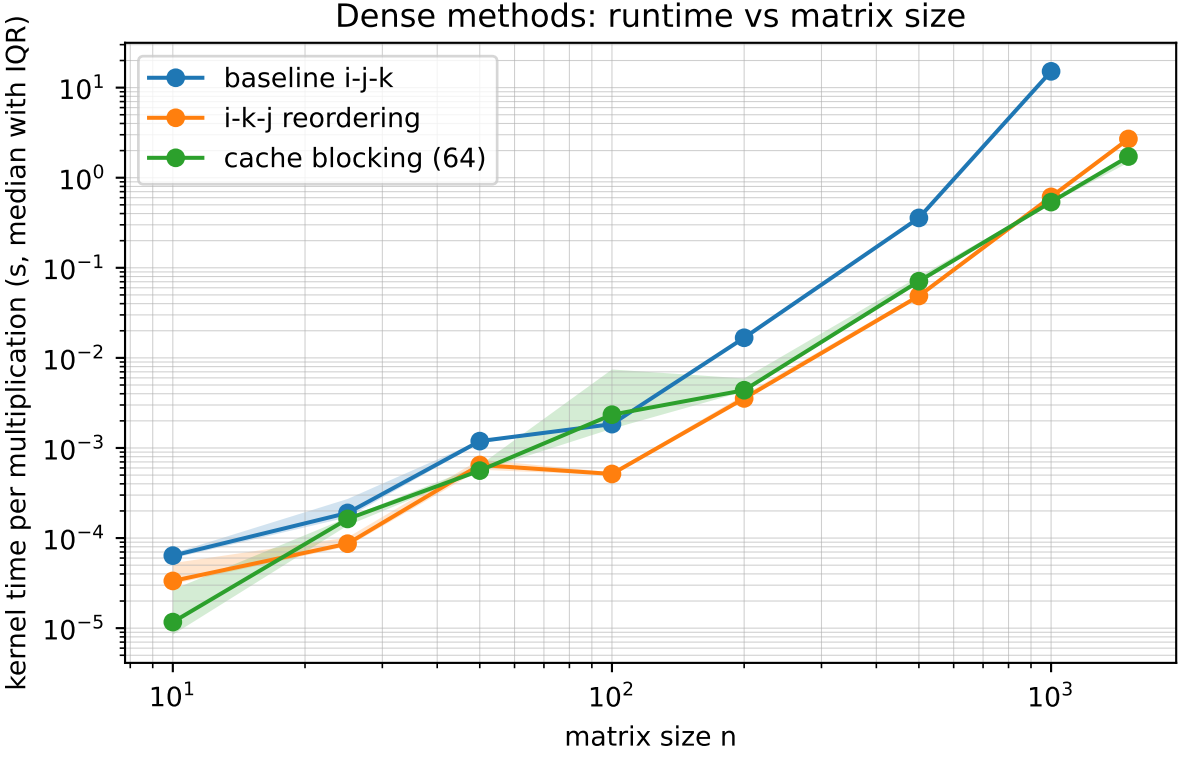
Dense methods: runtime vs matrix size
baseline i-j-k
i-k-j reordering
cache blocking (64)
kernel time per multiplication (s, median with IQR)
matrix size n
$10^{1}$
$10^{0}$
$10^{-1}$
$10^{-2}$
$10^{-3}$
$10^{-4}$
$10^{-5}$
$10^{1}$
$10^{2}$
$10^{3}$